

When FLOPs Mislead on PC GPUs: Comparing RTX 3050 and RTX 4060 Inference Energy

Aryan Shah

Texas Academy of Mathematics and Science
University of North Texas
Denton, TX, USA
aryanshah.2431@gmail.com

Romir Kadiam

Texas Academy of Mathematics and Science
University of North Texas
Denton, TX, USA
romirkadiam123@gmail.com

Abstract—This paper reports a Windows CUDA measurement artifact for neural-network inference energy on two commodity NVIDIA GPUs: an RTX 3050 4GB Laptop GPU and an RTX 4060 desktop GPU. For each GPU, five image-classification models are measured in a batch-1 stability run with ten 10-second windows per model and in a batch-size sweep over batch sizes 1, 4, 8, 16, and 32. The artifact retains raw `nvidia-smi` logs and reports GPU-only board energy, not CPU package energy or wall-plug energy. On the RTX 4060, latency explains most between-model energy variation ($R^2 = 0.963$), while FLOPs is weak ($R^2 = 0.038$). On the RTX 3050 Laptop GPU, latency remains more informative than FLOPs ($R^2 = 0.684$ versus 0.011) but the measurements are noisier. The batch sweeps show that best per-image energy depends on both model and GPU: the best batch size ranges from 8 to 32 on the RTX 4060 and from 16 to 32 on the RTX 3050. These results support a measurement-first PC-GPU benchmark methodology and show why FLOPs alone is not a reliable energy proxy.

Index Terms—GPU inference, energy measurement, Windows CUDA, RTX 3050, RTX 4060, PyTorch, benchmarking

I. Introduction

Model efficiency is often summarized with FLOPs, parameter count, or model size. Those quantities are useful architecture descriptors, but they do not directly measure the runtime behavior that determines deployment energy: operator scheduling, kernel launch overhead, memory movement, batching, and device utilization. This paper studies that gap with direct Windows CUDA measurements on two NVIDIA GPUs.

This paper is intentionally separate from the Apple-Silicon CPU-only study in the companion repository. The Apple paper uses macOS powermetrics; this PC paper uses `nvidia-smi` GPU-board telemetry. The absolute joules therefore have a different power boundary and should not be pooled with package-level or wall-plug measurements. The question here is local and systems-oriented: under a fixed Windows CUDA artifact, do architecture proxies such as FLOPs track measured GPU-only inference energy?

The current contributions are:

- A reproducible Windows CUDA workflow with raw `nvidia-smi` logs, environment metadata, randomized model order, and confidence intervals.
- Batch-1 energy measurements for five models on an RTX 3050 Laptop GPU and an RTX 4060 desktop GPU.
- A comparison showing that latency is a stronger predictor of GPU-only energy than FLOPs on both GPUs, though the strength differs by system.
- Batch-size sweeps showing that per-image energy changes by up to $20.1\times$ and that the best batch size is model- and GPU-dependent.
- A clear boundary statement: the artifact measures GPU board energy only, not CPU package energy or whole-system wall power.

II. Related Work and Motivation

Efficient model families such as MobileNet, EfficientNet, and residual networks make different tradeoffs among depthwise convolution, channel width, residual connectivity, and operator mix [1]–[3]. FLOPs and parameters are portable, but they do not capture kernel choice, memory traffic, launch overhead, or how well a given GPU is utilized. Prior model-efficiency studies have similarly shown that arithmetic count, parameter count, memory footprint, and measured speed can rank networks differently in practice [4]. At the hardware level, energy cost is also shaped by memory movement and data movement, not just arithmetic operations [5].

Systems benchmarking work therefore emphasizes measured deployment behavior. MLPerf Inference separates latency-oriented and throughput-oriented scenarios rather than assuming one number characterizes all deployments [6]. NVIDIA’s `nvidia-smi` and NVML tooling expose useful GPU telemetry, but that telemetry is device-level rather than whole-system energy [7]. This paper follows that measurement-first approach for a small, auditable PC-GPU artifact. The broader motivation is also environmental and operational: energy-aware machine learning work has argued that reporting compute and energy costs is necessary for transparent model evaluation [8].

TABLE I
Windows CUDA measurement environments.

Component	RTX 3050 system	RTX 4060 system
GPU	RTX 3050 4GB Laptop GPU	RTX 4060
CPU	Intel64 Model 186	Intel64 Model 191
Python / PyTorch	3.8.10 / 2.4.1+cu124	3.11.9 / 2.11.0+cu128
Power backend	nvidia-smi, GPU-only	nvidia-smi, GPU-only
Main protocol	10 windows/model, 10 s/window	10 windows/model, 10 s/window
Batch sweep	$B = 1, 4, 8, 16, 32$	$B = 1, 4, 8, 16, 32$

TABLE II
Measurement validity checks.

Risk	Check or mitigation
Startup overhead	discard 10 warmup inferences before each window
Order effects	randomize model order with recorded seed
Power parser failure	retain raw nvidia-smi output for every window
Missing samples	require nonzero parsed windows before energy claims
Boundary confusion	label all energy as GPU-only
Batch ambiguity	report batch size and per-image metrics separately

III. Measurement Scope

Table I summarizes the two systems. Both runs use CUDA, float32, batch size 1 for the main stability measurement, an input shape of $1 \times 3 \times 224 \times 224$, and the same five model definitions. The RTX 3050 run uses Python 3.8.10 and PyTorch 2.4.1+cu124; the RTX 4060 run uses Python 3.11.9 and PyTorch 2.11.0+cu128. Because software versions differ, the paper treats the systems as deployment measurements rather than as a controlled GPU microbenchmark.

The repository has four active data directories, one batch-1 directory and one batch-sweep directory for each GPU. Each contains the summary CSV, trial CSV, environment JSON, and raw power logs for the corresponding run.

IV. Power Logging

The measurement script first tries the structured nvidia-smi CSV query for instantaneous power. When that field is unavailable, the script falls back to parsing the raw nvidia-smi -q -d POWER text while preserving the raw log. Energy per inference is computed as

$$E_{\text{gpu,inf}} = \frac{\bar{P}_{\text{gpu}} \Delta t}{N}, \quad (1)$$

where $\bar{P}_{\text{gpu}}$ is mean sampled GPU board power, Δt is the measured window duration, and N is the number of completed inferences. For batch sweeps, per-image energy is

$$E_{\text{gpu,img}} = \frac{\bar{P}_{\text{gpu}} \Delta t}{NB}, \quad (2)$$

where B is batch size. These quantities are GPU-only and should not be read as CPU package or wall-plug energy.




Fig. 1. Batch-1 GPU-only energy versus latency. Latency is informative on both GPUs, but the RTX 3050 measurements are noisier.

V. Batch-1 Results

Table III gives the main batch-1 energy measurements. Every model has 10 parsed energy windows on both GPUs. The RTX 4060 is much faster, with $2.6\text{--}3.2\times$ lower latency than the RTX 3050 Laptop GPU, but it also draws substantially higher average board power. As a result, the RTX 4060 uses more GPU-board energy per batch-1 inference for every model in this artifact.

The key point is not that one GPU is universally “better” in absolute joules. The power boundary is GPU board power only, and the systems differ in driver stack, PyTorch version, and device class. The key point is that FLOPs does not determine the energy order on either system. ResNet-18 has the largest FLOP count but is the lowest-energy batch-1 model on both GPUs. EfficientNet-B0 has far fewer FLOPs than ResNet-18 but consumes more energy on both systems.

The intervals reveal a measurement difference between the systems. The RTX 4060 run has tight intervals because the board-power samples and inference counts are stable across windows. The RTX 3050 Laptop GPU run is usable, with all windows parsed, but it is visibly noisier. MobileNetV3-Small and EfficientNet-B0 show especially wide intervals, so their exact batch-1 energy values should be interpreted more cautiously than the RTX 4060 values. This does not change the qualitative ranking result: FLOPs remains a poor proxy, and ResNet-18 remains low energy despite having the largest FLOP count.

Table VI makes the power-latency tradeoff explicit. The RTX 4060 is between $2.58\times$ and $3.15\times$ faster at batch 1, but it uses $1.61\times$ to $2.27\times$ more GPU-board energy per inference. This is not a contradiction: the faster GPU operates at a much higher board-power level. It also means that an energy paper should avoid collapsing the result into a single winner. The energy-minimizing hardware under batch-1 GPU-board energy is not necessarily the latency-minimizing hardware, and neither is necessarily the best throughput hardware once batching is allowed.

TABLE III
Batch-1 Windows CUDA inference results. Energy is GPU-only from nvidia-smi power samples.

Model	FLOPs (G)	RTX 3050 Laptop			RTX 4060			4060/3050 Energy
		Energy (J)	Lat. (ms)	Power (W)	Energy (J)	Lat. (ms)	Power (W)	
MobileNetV3-S	0.113	0.107	13.55	8.43	0.236	4.31	54.91	2.21
MobileNetV2	0.602	0.109	11.31	10.21	0.248	4.38	56.67	2.27
ResNet-18	3.628	0.102	6.01	17.84	0.165	2.02	81.45	1.61
Tiny-ViT-5M	2.331	0.211	18.20	12.63	0.415	6.95	59.68	1.96
EffNet-B0	0.772	0.171	18.60	9.96	0.383	6.80	56.35	2.24

TABLE IV
Batch-1 95% confidence intervals for GPU-only energy. Values are joules per inference.

Model	RTX 3050	RTX 4060
MobileNetV3-S	[0.061, 0.152]	[0.232, 0.241]
MobileNetV2	[0.089, 0.129]	[0.245, 0.252]
ResNet-18	[0.084, 0.121]	[0.162, 0.167]
Tiny-ViT-5M	[0.167, 0.256]	[0.410, 0.419]
EffNet-B0	[0.123, 0.219]	[0.379, 0.388]

TABLE V
Univariate predictors of batch-1 GPU-only energy across five model means.

Predictor	RTX 3050	RTX 4060
Latency (ms/inference)	0.684	0.963
FLOPs (G)	0.011	0.038
Parameters (M)	0.190	0.020

VI. Batch-Size Sweeps

Batching changes the denominator of the energy calculation. A larger batch can draw more board power while still reducing per-image energy if throughput increases enough. Table VII reports the best observed batch size for each model and GPU. Both systems show large gains over batch 1, but the best batch size is not constant across models or GPUs.

Table VIII shows why the best-batch result should not be reduced to a single global recommendation. MobileNetV3-Small continues to benefit through batch 32 on both GPUs. MobileNetV2 reaches its best energy at batch 16 on the RTX 3050 but at batch 8 on the RTX 4060. Tiny-ViT-5M also peaks earlier on the RTX 4060 than on the RTX 3050. These differences are plausible because the two GPUs have different power envelopes and utilization behavior, and they reinforce the need to report batch size explicitly.

For throughput-oriented use, the RTX 4060 delivers much higher image rate at similar or slightly lower best per-image GPU energy. For example, MobileNetV3 reaches nearly the same best per-image energy on both GPUs, but the RTX 4060 achieves roughly $6.7\times$ the image rate at that point. This illustrates why energy, latency, and throughput should be reported together.




Fig. 2. Batch-1 GPU-only energy versus FLOPs. ResNet-18 has the largest FLOP count but is the lowest-energy model on both GPUs.

TABLE VI
Cross-GPU batch-1 deltas. Latency speedup is RTX 3050 latency divided by RTX 4060 latency.

Model	4060/3050 energy	4060 latency speedup
MobileNetV3-S	2.21	3.15
MobileNetV2	2.27	2.58
ResNet-18	1.61	2.97
Tiny-ViT-5M	1.96	2.62
EffNet-B0	2.24	2.73

VII. Discussion

The two-GPU artifact strengthens the original single-GPU result. FLOPs is weak on both systems, while latency is more predictive. However, the RTX 3050 result also shows that latency is not a universal law: its R^2 is lower and the batch-1 confidence intervals are wider. That difference is useful rather than embarrassing, because it shows that a publishable benchmark should expose system-specific behavior rather than forcing all devices into one proxy.

The RTX 4060 is consistently faster, but it draws much more board power. In batch-1 mode, the power increase is large enough that the RTX 4060 consumes more GPU-board energy per inference for these five models. Under batching, the RTX 4060 becomes more attractive for throughput-oriented operation because its image rate increases sharply while best per-image energy is similar to or lower than the RTX 3050. This is exactly why batch size is a first-order experimental condition, not an

TABLE VII

Best observed batch size by GPU-only energy per image. Gain compares batch size 1 with the best per-image energy in the sweep.

Model	RTX 3050 Laptop				RTX 4060			
	Best B	E/img	Images/s	Gain	Best B	E/img	Images/s	Gain
MobileNetV3-S	32	0.0130	1111	20.1	32	0.0127	7495	18.6
MobileNetV2	16	0.0523	280	3.8	8	0.0451	2012	5.4
ResNet-18	32	0.0479	311	2.9	16	0.0464	2114	3.3
Tiny-ViT-5M	16	0.1245	120	3.3	8	0.1039	958	3.9
EffNet-B0	16	0.0690	214	4.5	16	0.0668	1419	5.8

TABLE VIII

GPU-only energy per image across the batch sweeps. Values are joules per image.

Model	RTX 3050 Laptop					RTX 4060				
	B1	B4	B8	B16	B32	B1	B4	B8	B16	B32
MobileNetV3-S	0.2608	0.0824	0.0497	0.0318	0.0130	0.2354	0.0634	0.0352	0.0213	0.0127
MobileNetV2	0.2008	0.1003	0.0663	0.0523	0.0581	0.2438	0.0769	0.0451	0.0525	0.0560
ResNet-18	0.1403	0.0786	0.0589	0.0498	0.0479	0.1545	0.0582	0.0499	0.0464	0.0496
Tiny-ViT-5M	0.4102	0.1710	0.1424	0.1245	0.1284	0.4102	0.1333	0.1039	0.1076	0.1076
EffNet-B0	0.3092	0.1567	0.0955	0.0690	0.0725	0.3846	0.1184	0.0674	0.0668	0.0720

TABLE IX

Best-batch mean GPU board power. Values are watts at the energy-minimizing batch size in Table VII.

Model	RTX 3050	RTX 4060
MobileNetV3-S	14.4	95.0
MobileNetV2	14.6	90.7
ResNet-18	14.9	98.0
Tiny-ViT-5M	15.0	99.6
EffNet-B0	14.8	94.8




Fig. 3. Example batch-sweep curves. Batching helps both GPUs, but the energy curve and optimum are model- and device-dependent.

implementation detail.

The comparison also separates two meanings of “efficient.” For a latency-constrained interactive system, the RTX 4060 is preferable because it finishes each request much faster. For a strict GPU-board-energy-per-request metric at batch 1, the RTX 3050 Laptop GPU is lower energy in this artifact because its board power is much lower. For buffered throughput, the answer changes again: the RTX 4060 can process far more images per second at nearly the same best per-image energy. A paper that reports only FLOPs, or only batch-1 latency, would miss

TABLE X

Minimum metadata needed for each additional PC-GPU system.

Field	Reason
GPU and driver	power telemetry support differs by device and driver
CPU model	dispatch overhead is outside GPU energy but affects runtime
PyTorch/CUDA versions	kernel selection can change latency and power
Power backend	distinguishes GPU-only, CPU package, and wall power
Batch sizes	separates low-latency serving from throughput operation
Raw logs	allows parser failures and missing samples to be audited

those distinctions.

VIII. Protocol for Added Systems

The artifact is structured so another Windows CUDA system can be added without changing the paper workflow. A contributed system should include the complete output directory from the measurement script, not just the summary CSV. The minimum required files are `pc_inference_trials.csv`, `pc_inference_summary.csv`, `pc_measurement_environment.json`, and the complete `raw_power/` directory. The environment JSON is important because GPU model, driver behavior, CUDA version, PyTorch version, and Python version can all change measured latency and power.

For future systems, the batch-1 run should be collected first because it gives the cleanest predictor comparison. The batch sweep should then be collected with the same set $B \in \{1, 4, 8, 16, 32\}$ when memory permits. Runs with zero parsed power samples should be retained as latency data but excluded from energy claims until the parser or power backend is fixed.

TABLE XI
Windowed measurement loop used by the artifact.

Step	Action
1	Load the five model definitions and fix input shape, device, dtype, and seed.
2	Randomize model order for each repeat to reduce ordering effects.
3	Run warmup inferences before each measured window.
4	Start nvidia-smi logging and repeatedly execute forward passes for the target window duration.
5	Store raw power text, completed inference count, elapsed time, latency, and parsed sample count for the trial.
6	Aggregate trial rows into per-model summary rows with means, standard deviations, standard errors, and 95% confidence intervals.

IX. Measurement Implementation

The measurement script uses a windowed workload rather than timing one inference at a time. A single inference is too short for robust power sampling, especially on Windows where nvidia-smi reports power at a coarser cadence than CUDA kernel execution. Each trial therefore runs the model repeatedly for a fixed wall-clock window, counts completed inferences, samples GPU power in parallel, and divides total estimated GPU energy by the number of completed inferences.

This design has two practical advantages. First, it keeps the energy value auditable: if a summary row looks wrong, the corresponding trial rows and raw power logs can be inspected directly. Second, it makes failed telemetry visible instead of silently producing zeros. The paper only uses rows with nonzero parsed energy windows for energy claims. Latency values from failed power runs can still be useful, but they should not be mixed with direct energy results.

The batch sweep reuses the same loop but changes the input batch dimension. The summary CSV reports per-forward values, while the trial CSV also records per-image latency and energy. This distinction is necessary because a batch-32 forward pass performs 32 image inferences. In the paper, batch comparisons use per-image values; batch-1 predictor comparisons use the longer batch-1 stability runs.

X. Measurement Quality

The RTX 3050 Laptop GPU run is valid but visibly noisier than the RTX 4060 run. Table XII reports relative confidence-interval width, computed as $(CI_{high} - CI_{low})/\bar{E}$. This quantity is useful because it makes uncertainty comparable across models with different absolute energy values.

This table is important for interpretation. The RTX 3050 data should be used for broad patterns: FLOPs is weak, latency is more informative, and batching changes per-image energy substantially. It should not be overread as precise ranking evidence when confidence intervals overlap. The RTX 4060 data support more precise model-to-model comparisons because every relative interval is

TABLE XII
Relative width of the 95% energy confidence interval. Lower is better.

Model	RTX 3050	RTX 4060
MobileNetV3-S	85.0%	3.4%
MobileNetV2	36.4%	2.9%
ResNet-18	35.6%	2.7%
Tiny-ViT-5M	42.3%	2.1%
EffNet-B0	56.1%	2.4%

below 4%. A stronger future artifact would repeat the RTX 3050 run, align the software stack with the RTX 4060 system, and compare nvidia-smi against NVML or wall-meter measurements.

XI. Reproducibility and Validity

Each run includes a trial CSV, summary CSV, environment JSON, and raw power log directory. The paper uses a summary value only when all model windows have nonzero parsed power samples. The RTX 3050 and RTX 4060 batch-1 runs both pass that check with 10 parsed windows per model. The raw logs are retained so that parser failures or driver-specific power-field differences can be audited.

There are still important limits. The artifact measures GPU board energy, not whole-PC wall power. CPU dispatch, host memory, storage, display, cooling, and power-supply losses are outside the boundary. The two systems also differ in software stack, so the results should be read as deployment measurements, not as a pure architectural comparison of GPU silicon. Finally, the five-model regressions are descriptive; broader claims require more GPUs, repeated software versions, and ideally an external power meter for cross-checking.

XII. Artifact Availability

The organized repository stores the current paper in paper/. The RTX 3050 data are under data/rtx3050_batch1_main/ and data/rtx3050_batch_sweep/. The RTX 4060 data are under data/rtx4060_batch1_main/ and data/rtx4060_batch_sweep/. The original submitted RTX 3050 bundle is retained as aryan.zip. Each active directory includes the trial CSV, summary CSV, environment JSON, and raw power logs.

XIII. Conclusion

Across two Windows CUDA systems, FLOPs is a poor proxy for GPU-only inference energy. Latency is more informative on both GPUs, but the strength of that relationship is system-dependent. Batch size also changes the energy picture: best per-image energy varies by model and GPU, and the RTX 4060’s higher board power is offset by much higher throughput in batched operation. The result is a more credible PC-GPU paper than a single-device study because it shows both a common pattern and a device-specific difference.

References

- [1] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in Proc. IEEE Conference on Computer Vision and Pattern Recognition, 2016.
- [2] M. Sandler, A. Howard, M. Zhu, A. Zhmoginov, and L.-C. Chen, “MobileNetV2: Inverted residuals and linear bottlenecks,” in Proc. IEEE Conference on Computer Vision and Pattern Recognition, 2018.
- [3] M. Tan and Q. Le, “EfficientNet: Rethinking model scaling for convolutional neural networks,” in Proc. International Conference on Machine Learning, 2019.
- [4] A. Canziani, A. Paszke, and E. Culurciello, “An analysis of deep neural network models for practical applications,” arXiv:1605.07678, 2016.
- [5] M. Horowitz, “1.1 Computing’s energy problem (and what we can do about it),” in Proc. IEEE International Solid-State Circuits Conference, 2014.
- [6] V. J. Reddi et al., “MLPerf Inference Benchmark,” in Proc. ACM/IEEE International Symposium on Computer Architecture, 2020.
- [7] NVIDIA, “NVIDIA System Management Interface,” available: <https://developer.nvidia.com/system-management-interface>.
- [8] E. Strubell, A. Ganesh, and A. McCallum, “Energy and policy considerations for deep learning in NLP,” in Proc. ACL, 2019.